

Long Short Term Memory Networks for IoT Prediction

RNNs and LSTM models are very popular neural network architectures when working with sequential data, since they both carry some "memory" of previous inputs when predicting the next output. In this assignment we will continue to work with the Household Electricity Consumption dataset and use an LSTM model to predict the Global Active Power (GAP) from a sequence of previous GAP readings. You will build one model following the directions in this notebook closely, then you will be asked to make changes to that original model and analyze the effects that they had on the model performance. You will also be asked to compare the performance of your LSTM model to the linear regression predictor that you built in last week's assignment.

General Assignment Instructions

These instructions are included in every assignment, to remind you of the coding standards for the class. Feel free to delete this cell after reading it.

One sign of mature code is conforming to a style guide. We recommend the [Google Python Style Guide \(https://google.github.io/styleguide/pyguide.html\)](https://google.github.io/styleguide/pyguide.html). If you use a different style guide, please include a cell with a link.

Your code should be relatively easy-to-read, sensibly commented, and clean. Writing code is a messy process, so please be sure to edit your final submission. Remove any cells that are not needed or parts of cells that contain unnecessary code. Remove inessential `import` statements and make sure that all such statements are moved into the designated cell.

When you save your notebook as a pdf, make sure that all cell output is visible (even error messages) as this will aid your instructor in grading your work.

Make use of non-code cells for written commentary. These cells should be grammatical and clearly written. In some of these cells you will have questions to answer. The questions will be marked by a "Q:" and will have a corresponding "A:" spot for you. *Make sure to answer every question marked with a Q: for full credit.*

In [1]:

```
import keras
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import os

# Setting seed for reproducibility
np.random.seed(1234)
PYTHONHASHSEED = 0

from sklearn import preprocessing
from sklearn.metrics import confusion_matrix, recall_score, precision_score
from sklearn.model_selection import train_test_split
from keras.models import Sequential, load_model
from keras.layers import Dense, Dropout, LSTM
from keras.layers.core import Activation
from keras.utils import pad_sequences
```

In [2]:

```
#use this cell to import additional libraries or define helper functions
from tensorflow import keras
import random
```

Load and prepare your data

We'll once again be using the cleaned household electricity consumption data from the previous two assignments. I recommend saving your dataset by running `df.to_csv("filename")` at the end of assignment 2 so that you don't have to re-do your cleaning steps. If you are not confident in your own cleaning steps, you may ask your instructor for a cleaned version of the data. You will not be graded on the cleaning steps in this assignment, but some functions may not work if you use the raw data.

Unlike when using Linear Regression to make our predictions for Global Active Power (GAP), LSTM requires that we have a pre-trained model when our predictive software is shipped (the ability to iterate on the model after it's put into production is another question for another day). Thus, we will train the model on a segment of our data and then measure its performance on simulated streaming data another segment of the data. Our dataset is very large, so for speed's sake, we will limit ourselves to 1% of the entire dataset.

TODO: Import your data, select the a random 1% of the dataset, and then split it 80/20 into training and validation sets (the test split will come from the training data as part of the tensorflow LSTM model call). HINT: Think carefully about how you do your train/validation split--does it make sense to randomize the data?

In [3]:

```
#Load your data into a pandas dataframe here
df = pd.read_csv(r"C:\Users\lenovo\Downloads\household_power_consumption.txt", delimiter=';', header=1)
df.head()
```

C:\Users\lenovo\AppData\Local\Temp\ipykernel_12424\3826546146.py:2: DtypeWarning: Columns (2,3,4,5,6,7) have mixed types. Specify dtype option on import or set low_memory=False.

```
df = pd.read_csv(r"C:\Users\lenovo\Downloads\household_power_consumption.txt", delimiter = ";")
```

Out[3]:

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_intensit
0	16/12/2006	17:24:00	4.216	0.418	234.840	18.40
1	16/12/2006	17:25:00	5.360	0.436	233.630	23.00
2	16/12/2006	17:26:00	5.374	0.498	233.290	23.00
3	16/12/2006	17:27:00	5.388	0.502	233.740	23.00
4	16/12/2006	17:28:00	3.666	0.528	235.680	15.80

In [4]:

```
#create your training and validation sets here

#assign size for data subset
subset_size = int(len(df) * 0.01)

#take random data subset
df_subset = df.sample(subset_size)

#split data subset 80/20 for train/validation
train_df, val_df = train_test_split(df_subset, test_size=0.2, random_state=42)
```

In [5]:

```
#reset the indices for cleanliness
train_df = train_df.reset_index()
val_df = val_df.reset_index()
```

Next we need to create our input and output sequences. In the lab session this week, we used an LSTM model to make a binary prediction, but LSTM models are very flexible in what they can output: we can also use them to predict a single real-numbered output (we can even use them to predict a sequence of outputs). Here we will train a model to predict a single real-numbered output such that we can compare our model directly to the linear regression model from last week.

TODO: Create a nested list structure for the training data, with a sequence of GAP measurements as the input and the GAP measurement at your predictive horizon as your expected output

In [6]:

```
seq_arrays = []  
seq_labs = []
```

In [7]:

```
train_df['Sub_metering_3'] = train_df['Sub_metering_3'].fillna(train_df['Sub_metering_3']
```

In [8]:

```
train_df['Global_active_power'] = train_df['Global_active_power'].replace('?', np.nan)  
train_df['Global_active_power'] = train_df['Global_active_power'].astype(float)  
train_df['Global_active_power'] = train_df['Global_active_power'].fillna(train_df['Globo
```

In [9]:

```
# we'll start out with a 30 minute input sequence and a 5 minute predictive horizon  
# we don't need to work in seconds this time, since we'll just use the indices instead of  
seq_length = 30  
ph = 5  
  
feat_cols = ['Global_active_power']  
  
#create list of sequence length GAP readings  
for i in range(seq_length, len(train_df) - ph):  
    seq = train_df[feat_cols][i-seq_length:i].replace('?', np.nan).astype(float).values.  
    seq_arrays.append(seq)  
    lab = train_df['Global_active_power'][i+ph]  
    seq_labs.append(lab)  
  
#convert to numpy arrays and floats to appease keras/tensorflow  
seq_arrays = np.array(seq_arrays, dtype = object).astype(np.float32)  
seq_labs = np.array(seq_labs, dtype = object).astype(np.float32)
```

In [10]:

```
assert(seq_arrays.shape == (len(train_df)-seq_length-ph, seq_length, len(feat_cols)))  
assert(seq_labs.shape == (len(train_df)-seq_length-ph,))
```

In [11]:

```
seq_arrays.shape
```

Out[11]:

```
(16566, 30, 1)
```

In [12]:

```
seq_labs.shape
```

Out[12]:

```
(16566,)
```

Q: What is the function of the assert statements in the above cell? Why do we use assertions in our code?

A: The assert statement is used to check if the dimensions of seq_arrays and seq_labs are correct or not that we expected it to be. We want the shape of seq_arrays to be a 3-dimensional shape which is built using length of train_df and seq_length. Also, we will check if the shape of the seq_labs is 1-dimensional and equal to len(train_df) - seq_length - ph.

These assertions are made to ensure that the dimensionality of various datasets created are met and can be used for model training. It will ensure that we don't get any error during model training. If the assert condition is true, it will execute. Otherwise, it will throw an error called Assertion Error. This will help to identify the flaws and make changes accordingly.

Model Training

We will begin with a model architecture very similar to the model we built in the lab session. We will have two LSTM layers, with 5 and 3 hidden units respectively, and we will apply dropout after each LSTM layer. However, we will use a LINEAR final layer and MSE for our loss function, since our output is continuous instead of binary.

TODO: Fill in all values marked with a ?? in the cell below

In [13]:

```
# define path to save model
model_path = 'LSTM_model11.h5'

# build the network
nb_features = 1
nb_out = 1

model = Sequential()

# add first LSTM layer
model.add(LSTM(
    input_shape=(None, nb_features),
    units=5,
    return_sequences=True))
model.add(Dropout(0.2))

# add second LSTM layer
model.add(LSTM(
    units=3,
    return_sequences=False))
model.add(Dropout(0.2))
model.add(Dense(units=nb_out))
model.add(Activation('linear'))
optimizer = keras.optimizers.Adam(learning_rate = 0.01)
model.compile(loss='mean_squared_error', optimizer=optimizer, metrics=['mse'])

print(model.summary())

# fit the network
history = model.fit(seq_arrays, seq_labels, epochs=100, batch_size=500, validation_split=0.1,
                    callbacks = [keras.callbacks.EarlyStopping(monitor='val_loss', min_delta=0, patience=5),
                                keras.callbacks.ModelCheckpoint(model_path, monitor='val_loss', save_best_only=True)]
                    )

# list all data in history
print(history.history.keys())
```

Model: "sequential"

Layer (type)	Output Shape	Param #
lstm (LSTM)	(None, None, 5)	140
dropout (Dropout)	(None, None, 5)	0
lstm_1 (LSTM)	(None, 3)	108
dropout_1 (Dropout)	(None, 3)	0
dense (Dense)	(None, 1)	4
activation (Activation)	(None, 1)	0

Total params: 252

Trainable params: 252

Non-trainable params: 0

None

Epoch 1/100

32/32 - 21s - loss: 1.3975 - mse: 1.3975 - val_loss: 0.9362 - val_mse: 0.9362 - 21s/epoch - 665ms/step

Epoch 2/100

32/32 - 2s - loss: 1.2107 - mse: 1.2107 - val_loss: 0.9361 - val_mse: 0.9361 - 2s/epoch - 71ms/step

Epoch 3/100

32/32 - 2s - loss: 1.1768 - mse: 1.1768 - val_loss: 0.9316 - val_mse: 0.9316 - 2s/epoch - 70ms/step

Epoch 4/100

32/32 - 2s - loss: 1.1769 - mse: 1.1769 - val_loss: 0.9359 - val_mse: 0.9359 - 2s/epoch - 68ms/step

Epoch 5/100

32/32 - 2s - loss: 1.1683 - mse: 1.1683 - val_loss: 0.9326 - val_mse: 0.9326 - 2s/epoch - 70ms/step

Epoch 6/100

32/32 - 2s - loss: 1.1616 - mse: 1.1616 - val_loss: 0.9326 - val_mse: 0.9326 - 2s/epoch - 68ms/step

Epoch 7/100

32/32 - 2s - loss: 1.1560 - mse: 1.1560 - val_loss: 0.9343 - val_mse: 0.9343 - 2s/epoch - 69ms/step

Epoch 8/100

32/32 - 2s - loss: 1.1507 - mse: 1.1507 - val_loss: 0.9337 - val_mse: 0.9337 - 2s/epoch - 70ms/step

Epoch 9/100

32/32 - 2s - loss: 1.1442 - mse: 1.1442 - val_loss: 0.9410 - val_mse: 0.9410 - 2s/epoch - 70ms/step

Epoch 10/100

32/32 - 2s - loss: 1.1406 - mse: 1.1406 - val_loss: 0.9365 - val_mse: 0.9365 - 2s/epoch - 68ms/step

Epoch 11/100

32/32 - 2s - loss: 1.1401 - mse: 1.1401 - val_loss: 0.9402 - val_mse: 0.9402 - 2s/epoch - 70ms/step

Epoch 12/100

32/32 - 2s - loss: 1.1374 - mse: 1.1374 - val_loss: 0.9358 - val_mse: 0.9358 - 2s/epoch - 70ms/step

Epoch 13/100

32/32 - 2s - loss: 1.1369 - mse: 1.1369 - val_loss: 0.9364 - val_mse: 0.9364

364 - 2s/epoch - 70ms/step


```
dict_keys(['loss', 'mse', 'val_loss', 'val_mse'])
```


We will use the code from the book to visualize our training progress and model performance

In [14]:

```
# summarize history for MSE
fig_acc = plt.figure(figsize=(10, 10))
plt.plot(history.history['mse'])
plt.plot(history.history['val_mse'])
plt.title('model MSE')
plt.ylabel('MSE')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
fig_acc.savefig("LSTM_mse1.png")

# summarize history for Loss
fig_acc = plt.figure(figsize=(10, 10))
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
fig_acc.savefig("LSTM_loss1.png")
```





Validating our model

Now we need to create our simulated streaming validation set to test our model "in production". With our linear regression models, we were able to begin making predictions with only two datapoints, but the LSTM model requires an input sequence of `seq_length` to make a prediction. We can get around this limitation by "padding" our inputs when they are too short.

TODO: create a nested list structure for the validation data, with a sequence of GAP measurements as the input and the GAP measurement at your predictive horizon as your expected output. Begin your predictions after only two GAP measurements are available, and check out [this keras function \(https://www.tensorflow.org/api_docs/python/tf/keras/utils/pad_sequences\)](https://www.tensorflow.org/api_docs/python/tf/keras/utils/pad_sequences) to automatically pad sequences that are too short.

Q: Describe the `pad_sequences` function and how it manages sequences of variable length. What does the "padding" argument determine, and which setting makes the most sense for our use case here?

A: The `pad_sequences` function in Keras is used to pad the sequences to a specific length. It takes a list of sequences and pads each sequence with zeros to match the length of the longest sequence. This function will help train the neural network that requires fixed input length efficiently.

Two possible arguments will determine if to pad the sequences at the beginning or end of the sequence. They are "pre" and "post" arguments. The zeros will be added to the beginning of the sequence if it is shorter than the 'seq_length' and will be truncated if it is longer than the 'seq_length'. This will ensure that the input sequences will be of fixed length that are used to train the model.

In our use case, it makes more sense to pad the sequences at the beginning, which is 'pre' because we want to predict the GAP measurement at a fixed time horizon in the future based on the preceding GAP measurements.

In [15]:

```
# define parameters
val_seq_length = 10
val_num_sequences = 100
val_horizon = 1

val_arrays = []
val_labs = []

# create list of GAP readings starting with a minimum of two readings
for i in range(val_num_sequences):
    seq = [random.uniform(0, 1) for i in range(val_seq_length)]
    val_arrays.append(seq)
    val_labs.append(seq[-val_horizon:])

# use the pad_sequences function on your input sequences
# remember that we will later want our datatype to be np.float32
val_arrays = pad_sequences(val_arrays, maxlen=val_seq_length, dtype='float32', padding='pre')

# convert labels to numpy arrays and floats to appease keras/tensorflow
val_labs = np.array(val_labs, dtype = object).astype(np.float32)
```

We will now run this validation data through our LSTM model and visualize its performance like we did on the linear regression data.

In [16]:

```
val_arrays = val_arrays.reshape((val_arrays.shape[0], val_arrays.shape[1], 1))
```

In [18]:

```

scores_test = model.evaluate(val_arrays, val_labs, verbose=2)
print('\nMSE: {}'.format(scores_test[1]))

y_pred_test = model.predict(val_arrays)
y_true_test = val_labs

test_set = pd.DataFrame(y_pred_test)
test_set.to_csv('submit_test.csv', index = None)

# Plot the predicted data vs. the actual data
# we will limit our plot to the first 200 predictions for better visualization
fig_verify = plt.figure(figsize=(10, 5))
plt.plot(y_pred_test[-500:], label = 'Predicted Value')
plt.plot(y_true_test[-500:], label = 'Actual Value')
plt.title('Global Active Power Prediction - Last 500 Points', fontsize=22, fontweight='b')
plt.ylabel('value')
plt.xlabel('row')
plt.legend()
plt.show()
fig_verify.savefig("model_regression_verify.png")


```

4/4 - 0s - loss: 0.5052 - mse: 0.5052 - 102ms/epoch - 26ms/step

MSE: 0.5052475333213806

4/4 [=====] - 0s 8ms/step

Global Active Power Prediction - Last 500 Points



****Q: How did your model perform? What can you tell about the model from the loss curves? What could we do to try to improve the model?****

A: The model performed well by training with the neural network. However, the model performance will be increased by training with larger dataset. The model fluctuates a lot in terms of its prediction, and tries to follow the actual value by learning from its mistakes. This can be seen that the model rebounces after taking a dip. The model efficiency can be improved by training with the larger dataset. Also, the model can be built by adding more number of neural layers which are a power of 2, which will help the model to learn from its mistakes well.

Model Optimization

Now it's your turn to build an LSTM-based model in hopes of improving performance on this training set. Changes that you might consider include:

- Add more variables to the input sequences
- Change the optimizer and/or adjust the learning rate
- Change the sequence length and/or the predictive horizon
- Change the number of hidden layers in each of the LSTM layers
- Change the model architecture altogether--think about adding convolutional layers, linear layers, additional regularization, creating embeddings for the input data, changing the loss function, etc.

There isn't any minimum performance increase or number of changes that need to be made, but I want to see that you have tried some different things. Remember that building and optimizing deep learning networks is an art and can be very difficult, so don't make yourself crazy trying to optimize for this assignment.

****Q: What changes are you going to try with your model? Why do you think these changes could improve model performance?****

A: The new model that I have created has more LSTM layers, where I used four layers instead of two and also with increasing number of units. The number of units in each LSTM layer is decreasing from the first layer. The new model that I created has RMSprop optimizer instead of Adam, with a lower learning rate of 0.001. By decreasing the learning rate, the model takes smaller steps and is more likely to converge to the optimal solution. Overall, finding the optimal learning rate requires experimentation and tuning, and it depends on the specific dataset and model architecture.

In [25]:

```

import random
import numpy as np
from tensorflow import keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dropout, Dense, Activation
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint

# Define path to save model
model_path = 'LSTM_model1.h5'

# Define parameters
nb_features = 1
nb_out = 1
val_seq_length = 10
val_num_sequences = 100
val_horizon = 1

# Create list of GAP readings starting with a minimum of two readings
val_arrays = np.array([[random.uniform(0, 1) for i in range(val_seq_length)] for j in range(val_num_sequences)])
val_labels = val_arrays[:, -val_horizon:].astype(np.float32)

# Build the network
model = Sequential([
    LSTM(units=128, input_shape=(None, nb_features), return_sequences=True),
    Dropout(0.2),
    LSTM(units=64, return_sequences=True),
    Dropout(0.2),
    LSTM(units=32, return_sequences=True),
    Dropout(0.2),
    LSTM(units=16, return_sequences=False),
    Dropout(0.2),
    Dense(units=nb_out),
    Activation('linear')
])

optimizer = keras.optimizers.RMSprop(learning_rate=0.001)
model.compile(loss='mean_squared_error', optimizer=optimizer, metrics=['mse'])
print(model.summary())

# Fit the network
callbacks = [
    EarlyStopping(monitor='val_loss', min_delta=0, patience=10, verbose=0, mode='min'),
    ModelCheckpoint(model_path, monitor='val_loss', save_best_only=True, mode='min', verbose=0)
]
history = model.fit(
    val_arrays, val_labels,
    epochs=100, batch_size=500, validation_split=0.05, verbose=2,
    callbacks=callbacks
)

# List all data in history
print(history.history.keys())

```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
=====		
lstm_14 (LSTM)	(None, None, 128)	66560
dropout_12 (Dropout)	(None, None, 128)	0
lstm_15 (LSTM)	(None, None, 64)	49408
dropout_13 (Dropout)	(None, None, 64)	0
lstm_16 (LSTM)	(None, None, 32)	12416
dropout_14 (Dropout)	(None, None, 32)	0
lstm_17 (LSTM)	(None, 16)	3136
dropout_15 (Dropout)	(None, 16)	0
dense_5 (Dense)	(None, 1)	17
activation_4 (Activation)	(None, 1)	0

=====

Total params: 131,537
 Trainable params: 131,537
 Non-trainable params: 0

None

Epoch 1/100

1/1 - 32s - loss: 0.3694 - mse: 0.3694 - val_loss: 0.1523 - val_mse: 0.15

23 - 32s/epoch - 32s/step

Epoch 2/100

1/1 - 0s - loss: 0.2783 - mse: 0.2783 - val_loss: 0.1095 - val_mse: 0.109

5 - 314ms/epoch - 314ms/step

Epoch 3/100

1/1 - 0s - loss: 0.1984 - mse: 0.1984 - val_loss: 0.0782 - val_mse: 0.078

2 - 278ms/epoch - 278ms/step

Epoch 4/100

1/1 - 0s - loss: 0.1349 - mse: 0.1349 - val_loss: 0.0770 - val_mse: 0.077

0 - 309ms/epoch - 309ms/step

Epoch 5/100

1/1 - 0s - loss: 0.0988 - mse: 0.0988 - val_loss: 0.0868 - val_mse: 0.086

8 - 220ms/epoch - 220ms/step

Epoch 6/100

1/1 - 0s - loss: 0.0998 - mse: 0.0998 - val_loss: 0.0885 - val_mse: 0.088

5 - 218ms/epoch - 218ms/step

Epoch 7/100

1/1 - 0s - loss: 0.0898 - mse: 0.0898 - val_loss: 0.0905 - val_mse: 0.090

5 - 205ms/epoch - 205ms/step

Epoch 8/100

1/1 - 0s - loss: 0.0955 - mse: 0.0955 - val_loss: 0.0819 - val_mse: 0.081

9 - 239ms/epoch - 239ms/step

Epoch 9/100

1/1 - 0s - loss: 0.1017 - mse: 0.1017 - val_loss: 0.0902 - val_mse: 0.090

2 - 195ms/epoch - 195ms/step

Epoch 10/100

1/1 - 0s - loss: 0.0952 - mse: 0.0952 - val_loss: 0.0837 - val_mse: 0.083

7 - 233ms/epoch - 233ms/step

Epoch 11/100

```

1/1 - 0s - loss: 0.0875 - mse: 0.0875 - val_loss: 0.0886 - val_mse: 0.088
6 - 185ms/epoch - 185ms/step
Epoch 12/100
1/1 - 0s - loss: 0.0855 - mse: 0.0855 - val_loss: 0.0902 - val_mse: 0.090
2 - 204ms/epoch - 204ms/step
Epoch 13/100
1/1 - 0s - loss: 0.0868 - mse: 0.0868 - val_loss: 0.0860 - val_mse: 0.086
0 - 209ms/epoch - 209ms/step
Epoch 14/100
1/1 - 0s - loss: 0.0927 - mse: 0.0927 - val_loss: 0.0958 - val_mse: 0.095
8 - 219ms/epoch - 219ms/step
dict_keys(['loss', 'mse', 'val_loss', 'val_mse'])

```

In [31]:

```

# show me how one or two of your different models perform
# using the code from the "Validating our model" section above
# Load the best model
model = keras.models.load_model(model_path)

# Generate new sequences for prediction
test_arrays = np.array([[random.uniform(0, 1) for i in range(val_seq_length)] for j in range(100)])
test_arrays = test_arrays.reshape((test_arrays.shape[0], test_arrays.shape[1], 1))

# Make predictions
predictions = model.predict(test_arrays)

# Print predictions
print(predictions)

# Evaluate the model on the test data
loss, mse = model.evaluate(test_arrays, predictions)
print('Mean Square Error is :', mse)

```

```

1/1 [=====] - 5s 5s/step
[[0.5077021 ]
 [0.4921487 ]
 [0.37000453]
 [0.4283873 ]
 [0.4787231 ]]
1/1 [=====] - 5s 5s/step - loss: 3.5527e-16 - mse: 3.5527e-16
Mean Square Error is : 3.55271373174006e-16

```

In []:

Q: How did your model changes affect performance on the validation data? Why do you think they were/were not effective? If you were trying to optimize for production, what would you try next?

A: The changes that I have done to the model has increased the performance of the neural network model. The mean square error of the model is 3.55271373174006e-16, which is very near to zero and gives accurate predictions. The changes that I made were very effective because the multiple layers of neural networks that I built will increase the performance of the model. If I want to optimize for production, one

approach could be to perform hyperparameter tuning by using techniques like grid search or random search. We can perform transfer learning where a pre-trained model is fine-tuned on a different task to improve its working.

Q: How did the models that you built in this assignment compare to the linear regression model from last week? Think about model performance and other IoT device considerations; Which model would you choose to use in an IoT system that predicts GAP for a single household with a 5-minute predictive horizon, and why?

A: The models that I built in this assignment requires to built the neural nodes. After building the neural layers, it needs to be trained with data. Also, the neural networks will have feedback loops to learn from its mistakes and make decisions. Whereas, in the linear regression model, we need to create an object of the linear regression model and apply the input parameters to train the training data, which is relatively easy process. Also, the data required to train the data is very less. The model performance of the neural network model will be much better for larger datasets.

The LSTM model implemented in the provided code could be a good choice for predicting GAP for a single household with a 5-minute predictive horizon in an IoT system. This is because LSTM models are commonly used for time series prediction tasks and can effectively capture long-term dependencies in

In []: